

CS 795 Assignment 4 - SGD Optimizer

April 7, 2022

```
[1]: import torch
```

```
[2]: # Device configuration
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
device
```

```
[2]: device(type='cuda')
```

```
[3]: from torchvision import datasets
from torchvision.transforms import ToTensor
train_data = datasets.MNIST(
    root = 'data',
    train = True,
    transform = ToTensor(),
    download = True,
)
test_data = datasets.MNIST(
    root = 'data',
    train = False,
    transform = ToTensor()
)
```

```
[4]: print(train_data)
```

```
Dataset MNIST
  Number of datapoints: 60000
  Root location: data
  Split: Train
  StandardTransform
Transform: ToTensor()
```

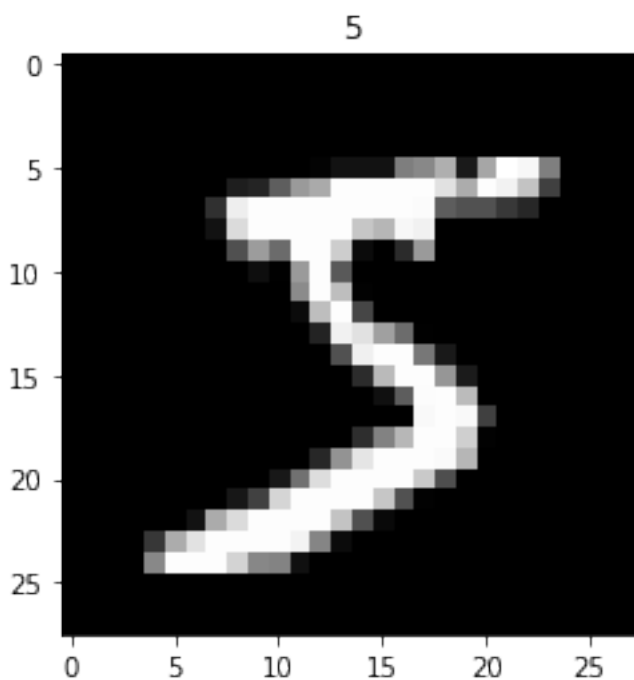
```
[5]: print(test_data)
```

```
Dataset MNIST
  Number of datapoints: 10000
  Root location: data
  Split: Test
  StandardTransform
Transform: ToTensor()
```


```
[6]: print(train_data.data.size())
```

```
torch.Size([60000, 28, 28])
```

```
[7]: import matplotlib.pyplot as plt
plt.imshow(train_data.data[0], cmap='gray')
plt.title('%i' % train_data.targets[0])
plt.show()
```



```
[8]: figure = plt.figure(figsize=(10, 8))
cols, rows = 5, 5
for i in range(1, cols * rows + 1):
    sample_idx = torch.randint(len(train_data), size=(1,)).item()
    img, label = train_data[sample_idx]
    figure.add_subplot(rows, cols, i)
    plt.title(label)
    plt.axis("off")
    plt.imshow(img.squeeze(), cmap="gray")
plt.show()
```



```
[9]: from torch.utils.data import DataLoader
loaders = {
    'train' : torch.utils.data.DataLoader(train_data,
                                          batch_size=100,
                                          shuffle=True,
                                          num_workers=1),

    'test'  : torch.utils.data.DataLoader(test_data,
                                          batch_size=100,
                                          shuffle=True,
                                          num_workers=1),
}
loaders
```

```
[9]: {'train': <torch.utils.data.dataloader.DataLoader at 0x7fe893e55be0>,
      'test': <torch.utils.data.dataloader.DataLoader at 0x7fe893e55a30>}
```

```
[10]: import torch.nn as nn
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(
            nn.Conv2d(
                in_channels=1,
                out_channels=16,
                kernel_size=5,
                stride=1,
                padding=2,
            ),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2),
        )
        self.conv2 = nn.Sequential(
            nn.Conv2d(16, 32, 5, 1, 2),
            nn.ReLU(),
            nn.MaxPool2d(2),
        )
        # fully connected layer, output 10 classes
        self.out = nn.Linear(32 * 7 * 7, 10)
    def forward(self, x):
        x = self.conv1(x)
        x = self.conv2(x)
        # flatten the output of conv2 to (batch_size, 32 * 7 * 7)
        x = x.view(x.size(0), -1)
        output = self.out(x)
        return output, x    # return x for visualization
```

```
[11]: cnn = CNN()
print(cnn)
```

```
CNN(
  (conv1): Sequential(
    (0): Conv2d(1, 16, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceiling_mode=False)
  )
  (conv2): Sequential(
    (0): Conv2d(16, 32, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceiling_mode=False)
  )
  (out): Linear(in_features=1568, out_features=10, bias=True)
)
```

```
[12]: loss_func = nn.CrossEntropyLoss()
      loss_func
```

```
[12]: CrossEntropyLoss()
```

```
[13]: import math
      from torch.optim import Optimizer

      class SGD(Optimizer):

          def __init__(self, params, lr=.01, momentum=0, dampening=0,
                        weight_decay=0, nesterov=False):
              defaults = dict(lr=lr, momentum=momentum, dampening=dampening,
                              weight_decay=weight_decay, nesterov=nesterov)
              super(SGD, self).__init__(params, defaults)

          def __setstate__(self, state):
              super(SGD, self).__setstate__(state)
              for group in self.param_groups:
                  group.setdefault('nesterov', False)

          def step(self, closure=None):
              loss = None
              if closure is not None:
                  loss = closure()

              for group in self.param_groups:
                  weight_decay = group['weight_decay']
                  momentum = group['momentum']
                  dampening = group['dampening']
                  nesterov = group['nesterov']

                  for p in group['params']:
                      if p.grad is None:
                          continue
                      d_p = p.grad.data
                      if weight_decay != 0:
                          d_p.add_(weight_decay, p.data)
                      # Apply learning rate
                      d_p.mul_(group['lr'])
                      if momentum != 0:
                          param_state = self.state[p]
                          if 'momentum_buffer' not in param_state:
                              buf = param_state['momentum_buffer'] = torch.
→zeros_like(p.data)
                              buf.mul_(momentum).add_(d_p)
                      else:
```

```

        buf = param_state['momentum_buffer']
        buf.mul_(momentum).add_(1 - dampening, d_p)
    if nesterov:
        d_p = d_p.add(momentum, buf)
    else:
        d_p = buf

    p.data.add_(-1, d_p)

    return loss

```

```

[14]: from torch import optim
optimizer = SGD(cnn.parameters(), lr = 0.01)
optimizer

```

```

[14]: SGD (
  Parameter Group 0
    dampening: 0
    lr: 0.01
    momentum: 0
    nesterov: False
    weight_decay: 0
)

```

```

[15]: from torch.autograd import Variable
num_epochs = 30
history = []
def test():
    # Test the model
    cnn.eval()
    with torch.no_grad():
        correct = 0
        total = 0
        for images, labels in loaders['test']:
            test_output, last_layer = cnn(images)
            b_y = Variable(labels)
            loss = loss_func(test_output, b_y)
            pred_y = torch.max(test_output, 1)[1].data.squeeze()
            accuracy = (pred_y == labels).sum().item() / float(labels.size(0))
            pass
        print('Test Accuracy of the model on the 10000 test images: %.2f' %
→accuracy)
        print('Test Loss: {:.4f}'.format(loss.item()))

    return accuracy, loss.item()

```

```

def train(num_epochs, cnn, loaders):

    cnn.train()

    # Train the model
    total_step = len(loaders['train'])

    for epoch in range(num_epochs):
        for i, (images, labels) in enumerate(loaders['train']):

            # gives batch data, normalize x when iterate train_loader
            b_x = Variable(images)    # batch x
            b_y = Variable(labels)    # batch y
            output = cnn(b_x)[0]
            loss = loss_func(output, b_y)

            # clear gradients for this training step
            optimizer.zero_grad()

            # backpropagation, compute gradients
            loss.backward()
            # apply gradients
            optimizer.step()

            pred_y = torch.max(output, 1)[1].data.squeeze()
            accuracy = (pred_y == labels).sum().item() / float(labels.size(0))
            if (i+1) % 100 == 0:
                print ('Epoch [{}/{}], Step [{}/{}], Acc: {:.4f} Loss: {:.4f}'
                        .format(epoch + 1, num_epochs, i + 1, total_step,
                                ↪accuracy, loss.item()))
                pass

            history.append([(accuracy, loss.item()), test()])
        pass
    train(num_epochs, cnn, loaders)

```

/tmp/ipykernel_1141096/781053462.py:49: UserWarning: This overload of add_ is deprecated:

```
    add_(Number alpha, Tensor other)
```

Consider using one of the following signatures instead:

```
    add_(Tensor other, *, Number alpha) (Triggered internally at
    ../torch/csrc/autograd/python_arg_parser.cpp:1050.)
    p.data.add_(-1, d_p)
```

Epoch [1/30], Step [100/600], Acc: 0.5800 Loss: 2.1339

Epoch [1/30], Step [200/600], Acc: 0.7700 Loss: 1.1554

Epoch [1/30], Step [300/600], Acc: 0.8600 Loss: 0.5479

Epoch [1/30], Step [400/600], Acc: 0.8700 Loss: 0.4957

Epoch [1/30], Step [500/600], Acc: 0.9400 Loss: 0.2994
 Epoch [1/30], Step [600/600], Acc: 0.8600 Loss: 0.3850
 Test Accuracy of the model on the 10000 test images: 0.91
 Test Loss: {:.4f} 0.3080783784389496
 Epoch [2/30], Step [100/600], Acc: 0.9400 Loss: 0.1773
 Epoch [2/30], Step [200/600], Acc: 0.9400 Loss: 0.3130
 Epoch [2/30], Step [300/600], Acc: 0.9600 Loss: 0.1939
 Epoch [2/30], Step [400/600], Acc: 0.8900 Loss: 0.3680
 Epoch [2/30], Step [500/600], Acc: 0.9300 Loss: 0.2190
 Epoch [2/30], Step [600/600], Acc: 0.8900 Loss: 0.2287
 Test Accuracy of the model on the 10000 test images: 0.93
 Test Loss: {:.4f} 0.161906898021698
 Epoch [3/30], Step [100/600], Acc: 0.9400 Loss: 0.3019
 Epoch [3/30], Step [200/600], Acc: 0.9300 Loss: 0.2088
 Epoch [3/30], Step [300/600], Acc: 0.9400 Loss: 0.2661
 Epoch [3/30], Step [400/600], Acc: 0.9000 Loss: 0.1892
 Epoch [3/30], Step [500/600], Acc: 0.9900 Loss: 0.1028
 Epoch [3/30], Step [600/600], Acc: 0.9300 Loss: 0.2188
 Test Accuracy of the model on the 10000 test images: 0.96
 Test Loss: {:.4f} 0.19273878633975983
 Epoch [4/30], Step [100/600], Acc: 0.9600 Loss: 0.1795
 Epoch [4/30], Step [200/600], Acc: 0.9400 Loss: 0.1431
 Epoch [4/30], Step [300/600], Acc: 0.9600 Loss: 0.1107
 Epoch [4/30], Step [400/600], Acc: 0.9800 Loss: 0.0863
 Epoch [4/30], Step [500/600], Acc: 0.9600 Loss: 0.1264
 Epoch [4/30], Step [600/600], Acc: 0.9500 Loss: 0.2360
 Test Accuracy of the model on the 10000 test images: 0.97
 Test Loss: {:.4f} 0.08885536342859268
 Epoch [5/30], Step [100/600], Acc: 0.9500 Loss: 0.1313
 Epoch [5/30], Step [200/600], Acc: 0.9500 Loss: 0.1281
 Epoch [5/30], Step [300/600], Acc: 0.9800 Loss: 0.0794
 Epoch [5/30], Step [400/600], Acc: 0.9600 Loss: 0.1421
 Epoch [5/30], Step [500/600], Acc: 0.9600 Loss: 0.1243
 Epoch [5/30], Step [600/600], Acc: 0.9800 Loss: 0.0856
 Test Accuracy of the model on the 10000 test images: 0.98
 Test Loss: {:.4f} 0.08267923444509506
 Epoch [6/30], Step [100/600], Acc: 0.9400 Loss: 0.1503
 Epoch [6/30], Step [200/600], Acc: 0.9600 Loss: 0.0886
 Epoch [6/30], Step [300/600], Acc: 0.9700 Loss: 0.0655
 Epoch [6/30], Step [400/600], Acc: 0.9600 Loss: 0.1211
 Epoch [6/30], Step [500/600], Acc: 0.9800 Loss: 0.1072
 Epoch [6/30], Step [600/600], Acc: 0.9700 Loss: 0.1110
 Test Accuracy of the model on the 10000 test images: 0.98
 Test Loss: {:.4f} 0.1043318659067154
 Epoch [7/30], Step [100/600], Acc: 0.9800 Loss: 0.0562
 Epoch [7/30], Step [200/600], Acc: 1.0000 Loss: 0.0411
 Epoch [7/30], Step [300/600], Acc: 0.9800 Loss: 0.0708
 Epoch [7/30], Step [400/600], Acc: 1.0000 Loss: 0.0351

Epoch [7/30], Step [500/600], Acc: 0.9700 Loss: 0.1597
 Epoch [7/30], Step [600/600], Acc: 0.9700 Loss: 0.1258
 Test Accuracy of the model on the 10000 test images: 1.00
 Test Loss: {:.4f} 0.0709439367055893
 Epoch [8/30], Step [100/600], Acc: 0.9700 Loss: 0.1121
 Epoch [8/30], Step [200/600], Acc: 0.9800 Loss: 0.0769
 Epoch [8/30], Step [300/600], Acc: 0.9800 Loss: 0.0508
 Epoch [8/30], Step [400/600], Acc: 0.9800 Loss: 0.0929
 Epoch [8/30], Step [500/600], Acc: 0.9500 Loss: 0.1171
 Epoch [8/30], Step [600/600], Acc: 0.9900 Loss: 0.0429
 Test Accuracy of the model on the 10000 test images: 0.98
 Test Loss: {:.4f} 0.10592630505561829
 Epoch [9/30], Step [100/600], Acc: 0.9500 Loss: 0.1405
 Epoch [9/30], Step [200/600], Acc: 0.9700 Loss: 0.1236
 Epoch [9/30], Step [300/600], Acc: 0.9600 Loss: 0.2261
 Epoch [9/30], Step [400/600], Acc: 0.9800 Loss: 0.0765
 Epoch [9/30], Step [500/600], Acc: 0.9800 Loss: 0.0855
 Epoch [9/30], Step [600/600], Acc: 0.9800 Loss: 0.0783
 Test Accuracy of the model on the 10000 test images: 0.98
 Test Loss: {:.4f} 0.0874273031949997
 Epoch [10/30], Step [100/600], Acc: 0.9900 Loss: 0.0460
 Epoch [10/30], Step [200/600], Acc: 0.9900 Loss: 0.0421
 Epoch [10/30], Step [300/600], Acc: 0.9900 Loss: 0.0559
 Epoch [10/30], Step [400/600], Acc: 0.9700 Loss: 0.1195
 Epoch [10/30], Step [500/600], Acc: 0.9900 Loss: 0.0373
 Epoch [10/30], Step [600/600], Acc: 0.9900 Loss: 0.0278
 Test Accuracy of the model on the 10000 test images: 0.97
 Test Loss: {:.4f} 0.08896807581186295
 Epoch [11/30], Step [100/600], Acc: 0.9900 Loss: 0.0279
 Epoch [11/30], Step [200/600], Acc: 0.9800 Loss: 0.0478
 Epoch [11/30], Step [300/600], Acc: 1.0000 Loss: 0.0285
 Epoch [11/30], Step [400/600], Acc: 0.9600 Loss: 0.1064
 Epoch [11/30], Step [500/600], Acc: 0.9700 Loss: 0.0929
 Epoch [11/30], Step [600/600], Acc: 0.9500 Loss: 0.1744
 Test Accuracy of the model on the 10000 test images: 1.00
 Test Loss: {:.4f} 0.03506298363208771
 Epoch [12/30], Step [100/600], Acc: 0.9800 Loss: 0.0555
 Epoch [12/30], Step [200/600], Acc: 0.9800 Loss: 0.1073
 Epoch [12/30], Step [300/600], Acc: 0.9800 Loss: 0.0927
 Epoch [12/30], Step [400/600], Acc: 1.0000 Loss: 0.0141
 Epoch [12/30], Step [500/600], Acc: 1.0000 Loss: 0.0198
 Epoch [12/30], Step [600/600], Acc: 1.0000 Loss: 0.0127
 Test Accuracy of the model on the 10000 test images: 0.98
 Test Loss: {:.4f} 0.06938844919204712
 Epoch [13/30], Step [100/600], Acc: 0.9900 Loss: 0.0261
 Epoch [13/30], Step [200/600], Acc: 0.9900 Loss: 0.0385
 Epoch [13/30], Step [300/600], Acc: 1.0000 Loss: 0.0317
 Epoch [13/30], Step [400/600], Acc: 0.9900 Loss: 0.0337

Epoch [13/30], Step [500/600], Acc: 0.9600 Loss: 0.0775
Epoch [13/30], Step [600/600], Acc: 0.9900 Loss: 0.0528
Test Accuracy of the model on the 10000 test images: 1.00
Test Loss: {:.4f} 0.0140267014503479
Epoch [14/30], Step [100/600], Acc: 1.0000 Loss: 0.0151
Epoch [14/30], Step [200/600], Acc: 0.9600 Loss: 0.0748
Epoch [14/30], Step [300/600], Acc: 0.9600 Loss: 0.1147
Epoch [14/30], Step [400/600], Acc: 1.0000 Loss: 0.0111
Epoch [14/30], Step [500/600], Acc: 0.9300 Loss: 0.1434
Epoch [14/30], Step [600/600], Acc: 0.9800 Loss: 0.0479
Test Accuracy of the model on the 10000 test images: 0.99
Test Loss: {:.4f} 0.030371250584721565
Epoch [15/30], Step [100/600], Acc: 0.9700 Loss: 0.0537
Epoch [15/30], Step [200/600], Acc: 0.9700 Loss: 0.1515
Epoch [15/30], Step [300/600], Acc: 0.9800 Loss: 0.0334
Epoch [15/30], Step [400/600], Acc: 0.9800 Loss: 0.0587
Epoch [15/30], Step [500/600], Acc: 0.9800 Loss: 0.0477
Epoch [15/30], Step [600/600], Acc: 0.9900 Loss: 0.0428
Test Accuracy of the model on the 10000 test images: 0.96
Test Loss: {:.4f} 0.08501268178224564
Epoch [16/30], Step [100/600], Acc: 0.9900 Loss: 0.0298
Epoch [16/30], Step [200/600], Acc: 0.9800 Loss: 0.0458
Epoch [16/30], Step [300/600], Acc: 0.9800 Loss: 0.0404
Epoch [16/30], Step [400/600], Acc: 1.0000 Loss: 0.0174
Epoch [16/30], Step [500/600], Acc: 0.9900 Loss: 0.0244
Epoch [16/30], Step [600/600], Acc: 0.9900 Loss: 0.0486
Test Accuracy of the model on the 10000 test images: 0.99
Test Loss: {:.4f} 0.049576807767152786
Epoch [17/30], Step [100/600], Acc: 0.9700 Loss: 0.0515
Epoch [17/30], Step [200/600], Acc: 0.9800 Loss: 0.0764
Epoch [17/30], Step [300/600], Acc: 1.0000 Loss: 0.0134
Epoch [17/30], Step [400/600], Acc: 0.9900 Loss: 0.0307
Epoch [17/30], Step [500/600], Acc: 1.0000 Loss: 0.0172
Epoch [17/30], Step [600/600], Acc: 0.9800 Loss: 0.0804
Test Accuracy of the model on the 10000 test images: 0.94
Test Loss: {:.4f} 0.11730123311281204
Epoch [18/30], Step [100/600], Acc: 0.9700 Loss: 0.0901
Epoch [18/30], Step [200/600], Acc: 0.9900 Loss: 0.0448
Epoch [18/30], Step [300/600], Acc: 1.0000 Loss: 0.0284
Epoch [18/30], Step [400/600], Acc: 0.9800 Loss: 0.0834
Epoch [18/30], Step [500/600], Acc: 0.9700 Loss: 0.0990
Epoch [18/30], Step [600/600], Acc: 0.9900 Loss: 0.0504
Test Accuracy of the model on the 10000 test images: 0.99
Test Loss: {:.4f} 0.019461361691355705
Epoch [19/30], Step [100/600], Acc: 1.0000 Loss: 0.0160
Epoch [19/30], Step [200/600], Acc: 0.9900 Loss: 0.0474
Epoch [19/30], Step [300/600], Acc: 0.9800 Loss: 0.0542
Epoch [19/30], Step [400/600], Acc: 0.9900 Loss: 0.0368

Epoch [19/30], Step [500/600], Acc: 0.9900 Loss: 0.0244
 Epoch [19/30], Step [600/600], Acc: 1.0000 Loss: 0.0203
 Test Accuracy of the model on the 10000 test images: 0.99
 Test Loss: {:.4f} 0.018195874989032745
 Epoch [20/30], Step [100/600], Acc: 0.9800 Loss: 0.0245
 Epoch [20/30], Step [200/600], Acc: 0.9900 Loss: 0.0416
 Epoch [20/30], Step [300/600], Acc: 0.9900 Loss: 0.0235
 Epoch [20/30], Step [400/600], Acc: 0.9900 Loss: 0.0601
 Epoch [20/30], Step [500/600], Acc: 0.9800 Loss: 0.0509
 Epoch [20/30], Step [600/600], Acc: 1.0000 Loss: 0.0235
 Test Accuracy of the model on the 10000 test images: 0.98
 Test Loss: {:.4f} 0.09288153797388077
 Epoch [21/30], Step [100/600], Acc: 0.9800 Loss: 0.0360
 Epoch [21/30], Step [200/600], Acc: 0.9800 Loss: 0.0870
 Epoch [21/30], Step [300/600], Acc: 0.9900 Loss: 0.0426
 Epoch [21/30], Step [400/600], Acc: 0.9800 Loss: 0.0964
 Epoch [21/30], Step [500/600], Acc: 0.9700 Loss: 0.0587
 Epoch [21/30], Step [600/600], Acc: 0.9900 Loss: 0.0441
 Test Accuracy of the model on the 10000 test images: 0.98
 Test Loss: {:.4f} 0.04535512626171112
 Epoch [22/30], Step [100/600], Acc: 1.0000 Loss: 0.0187
 Epoch [22/30], Step [200/600], Acc: 0.9800 Loss: 0.0501
 Epoch [22/30], Step [300/600], Acc: 0.9900 Loss: 0.0309
 Epoch [22/30], Step [400/600], Acc: 0.9700 Loss: 0.0640
 Epoch [22/30], Step [500/600], Acc: 0.9900 Loss: 0.0204
 Epoch [22/30], Step [600/600], Acc: 0.9900 Loss: 0.0705
 Test Accuracy of the model on the 10000 test images: 1.00
 Test Loss: {:.4f} 0.015459936112165451
 Epoch [23/30], Step [100/600], Acc: 1.0000 Loss: 0.0152
 Epoch [23/30], Step [200/600], Acc: 1.0000 Loss: 0.0220
 Epoch [23/30], Step [300/600], Acc: 0.9800 Loss: 0.0498
 Epoch [23/30], Step [400/600], Acc: 0.9900 Loss: 0.0283
 Epoch [23/30], Step [500/600], Acc: 0.9900 Loss: 0.0305
 Epoch [23/30], Step [600/600], Acc: 0.9900 Loss: 0.0476
 Test Accuracy of the model on the 10000 test images: 1.00
 Test Loss: {:.4f} 0.013581983745098114
 Epoch [24/30], Step [100/600], Acc: 0.9700 Loss: 0.0452
 Epoch [24/30], Step [200/600], Acc: 1.0000 Loss: 0.0055
 Epoch [24/30], Step [300/600], Acc: 0.9800 Loss: 0.0543
 Epoch [24/30], Step [400/600], Acc: 0.9800 Loss: 0.0790
 Epoch [24/30], Step [500/600], Acc: 1.0000 Loss: 0.0105
 Epoch [24/30], Step [600/600], Acc: 0.9900 Loss: 0.0388
 Test Accuracy of the model on the 10000 test images: 0.99
 Test Loss: {:.4f} 0.029347749426960945
 Epoch [25/30], Step [100/600], Acc: 0.9800 Loss: 0.0335
 Epoch [25/30], Step [200/600], Acc: 0.9800 Loss: 0.0671
 Epoch [25/30], Step [300/600], Acc: 0.9900 Loss: 0.0293
 Epoch [25/30], Step [400/600], Acc: 1.0000 Loss: 0.0251

Epoch [25/30], Step [500/600], Acc: 0.9600 Loss: 0.0648
 Epoch [25/30], Step [600/600], Acc: 0.9800 Loss: 0.0426
 Test Accuracy of the model on the 10000 test images: 0.99
 Test Loss: {:.4f} 0.03627597913146019
 Epoch [26/30], Step [100/600], Acc: 0.9900 Loss: 0.0279
 Epoch [26/30], Step [200/600], Acc: 0.9900 Loss: 0.0336
 Epoch [26/30], Step [300/600], Acc: 0.9900 Loss: 0.0245
 Epoch [26/30], Step [400/600], Acc: 0.9900 Loss: 0.0485
 Epoch [26/30], Step [500/600], Acc: 1.0000 Loss: 0.0211
 Epoch [26/30], Step [600/600], Acc: 0.9900 Loss: 0.0292
 Test Accuracy of the model on the 10000 test images: 0.97
 Test Loss: {:.4f} 0.08091029524803162
 Epoch [27/30], Step [100/600], Acc: 0.9800 Loss: 0.0677
 Epoch [27/30], Step [200/600], Acc: 1.0000 Loss: 0.0204
 Epoch [27/30], Step [300/600], Acc: 1.0000 Loss: 0.0209
 Epoch [27/30], Step [400/600], Acc: 1.0000 Loss: 0.0152
 Epoch [27/30], Step [500/600], Acc: 0.9900 Loss: 0.0237
 Epoch [27/30], Step [600/600], Acc: 0.9900 Loss: 0.0202
 Test Accuracy of the model on the 10000 test images: 1.00
 Test Loss: {:.4f} 0.0191293153911829
 Epoch [28/30], Step [100/600], Acc: 0.9900 Loss: 0.0464
 Epoch [28/30], Step [200/600], Acc: 0.9800 Loss: 0.1030
 Epoch [28/30], Step [300/600], Acc: 1.0000 Loss: 0.0164
 Epoch [28/30], Step [400/600], Acc: 1.0000 Loss: 0.0022
 Epoch [28/30], Step [500/600], Acc: 0.9900 Loss: 0.0158
 Epoch [28/30], Step [600/600], Acc: 1.0000 Loss: 0.0365
 Test Accuracy of the model on the 10000 test images: 0.99
 Test Loss: {:.4f} 0.03427070006728172
 Epoch [29/30], Step [100/600], Acc: 1.0000 Loss: 0.0092
 Epoch [29/30], Step [200/600], Acc: 0.9600 Loss: 0.0785
 Epoch [29/30], Step [300/600], Acc: 0.9900 Loss: 0.0418
 Epoch [29/30], Step [400/600], Acc: 1.0000 Loss: 0.0149
 Epoch [29/30], Step [500/600], Acc: 0.9900 Loss: 0.0453
 Epoch [29/30], Step [600/600], Acc: 0.9800 Loss: 0.0504
 Test Accuracy of the model on the 10000 test images: 0.98
 Test Loss: {:.4f} 0.04289078339934349
 Epoch [30/30], Step [100/600], Acc: 0.9900 Loss: 0.0362
 Epoch [30/30], Step [200/600], Acc: 1.0000 Loss: 0.0067
 Epoch [30/30], Step [300/600], Acc: 1.0000 Loss: 0.0055
 Epoch [30/30], Step [400/600], Acc: 0.9900 Loss: 0.0373
 Epoch [30/30], Step [500/600], Acc: 0.9900 Loss: 0.0289
 Epoch [30/30], Step [600/600], Acc: 0.9900 Loss: 0.0216
 Test Accuracy of the model on the 10000 test images: 1.00
 Test Loss: {:.4f} 0.0037076130975037813

```

[16]: x = range(1, num_epochs+1)
      acc = [i[0][0] for i in history]
  
```


```

val_acc = [i[1][0] for i in history]
loss = [i[1][0] for i in history]
val_loss = [i[1][1] for i in history]
plt.subplot(1, 2, 1)
plt.plot(x, acc, label="Accuracy")
plt.plot(x, val_acc, label="Validation Accuracy")
plt.legend(loc='upper right')
plt.title("Training and validation Accuracy")

plt.subplot(1, 2, 2)
plt.plot(x, loss, label="Loss")
plt.plot(x, val_loss, label="Validation Loss")
plt.legend(loc='upper right')
plt.title("Training and validation Loss")

plt.show()

```



[]:

[]: