

CS795-1c-b

April 19, 2022

```
[1]: !pip install pytorch-ignite torchsummary
```

```
Requirement already satisfied: pytorch-ignite in  
/home/pankaj/anaconda3/lib/python3.9/site-packages (0.4.8)  
Requirement already satisfied: torchsummary in  
/home/pankaj/anaconda3/lib/python3.9/site-packages (1.5.1)  
Requirement already satisfied: torch<2,>=1.3 in  
/home/pankaj/anaconda3/lib/python3.9/site-packages (from pytorch-ignite)  
(1.10.0+cu113)  
Requirement already satisfied: typing-extensions in  
/home/pankaj/.local/lib/python3.9/site-packages (from torch<2,>=1.3->pytorch-  
ignite) (4.0.0)
```

```
[2]: import torch  
import torchvision  
import ignite  
  
print(*map(lambda m: ": ".join((m.__name__, m.__version__)), (torch, ↵  
↵torchvision, ignite)), sep="\n")
```

```
torch: 1.10.0+cu113  
torchvision: 0.11.1+cu113  
ignite: 0.4.8
```

0.1 Import Libraries

Note: torchsummary is an optional dependency here.

```
[3]: import os  
import logging  
import matplotlib.pyplot as plt  
  
import numpy as np  
  
from torchsummary import summary  
  
import torch  
import torch.nn as nn  
import torch.optim as optim
```

```

import torchvision.transforms as transforms
import torchvision.utils as vutils

from ignite.engine import Engine, Events
import ignite.distributed as idist

```

0.2 Reproducibility and logging details

```
[4]: ignite.utils.manual_seed(999)
```

```

[5]: ignite.utils.setup_logger(name="ignite.distributed.auto.auto_dataloader",
    ↪level=logging.WARNING)
ignite.utils.setup_logger(name="ignite.distributed.launcher.Parallel",
    ↪level=logging.WARNING)

```

```
[5]: <Logger ignite.distributed.launcher.Parallel (WARNING)>
```

```

[6]: import math
import torch
from torch.optim import Optimizer

required = object()

class Extragradient(Optimizer):

    def __init__(self, params, defaults):
        super(Extragradient, self).__init__(params, defaults)
        self.params_copy = []

    def update(self, p, group):
        raise NotImplementedError

    def extrapolation(self):
        # Check if a copy of the parameters was already made.
        is_empty = len(self.params_copy) == 0
        for group in self.param_groups:
            for p in group['params']:
                u = self.update(p, group)
                if is_empty:
                    # Save the current parameters for the update step.
                    # Several extrapolation step can be made before each update
                    # but only the parameters before the first extrapolation
                    ↪step
                    # are saved.

```

```

        self.params_copy.append(p.data.clone())
    if u is None:
        continue
    # Update the current parameters
    p.data.add_(u)

def step(self, closure=None):

    if len(self.params_copy) == 0:
        raise RuntimeError('Need to call extrapolation before calling step.
→')

    loss = None
    if closure is not None:
        loss = closure()

    i = -1
    for group in self.param_groups:
        for p in group['params']:
            i += 1
            u = self.update(p, group)
            if u is None:
                continue
            # Update the parameters saved during the extrapolation step
            p.data = self.params_copy[i].add_(u)

    # Free the old parameters
    self.params_copy = []
    return loss

class ExtraSGD(Extragradient):

    def __init__(self, params, lr=required, momentum=0, dampening=0,
                  weight_decay=0, nesterov=False):
        if lr is not required and lr < 0.0:
            raise ValueError("Invalid learning rate: {}".format(lr))
        if momentum < 0.0:
            raise ValueError("Invalid momentum value: {}".format(momentum))
        if weight_decay < 0.0:
            raise ValueError("Invalid weight_decay value: {}".
→format(weight_decay))

        defaults = dict(lr=lr, momentum=momentum, dampening=dampening,
                        weight_decay=weight_decay, nesterov=nesterov)
        if nesterov and (momentum <= 0 or dampening != 0):

```

```

        raise ValueError("Nesterov momentum requires a momentum and zero_
→dampening")
    super(ExtraSGD, self).__init__(params, defaults)

    def __setstate__(self, state):
        super(ExtraSGD, self).__setstate__(state)
        for group in self.param_groups:
            group.setdefault('nesterov', False)

    def update(self, p, group):
        weight_decay = group['weight_decay']
        momentum = group['momentum']
        dampening = group['dampening']
        nesterov = group['nesterov']

        if p.grad is None:
            return None
        d_p = p.grad.data
        if weight_decay != 0:
            d_p.add_(weight_decay, p.data)
        if momentum != 0:
            param_state = self.state[p]
            if 'momentum_buffer' not in param_state:
                buf = param_state['momentum_buffer'] = torch.zeros_like(p.data)
                buf.mul_(momentum).add_(d_p)
            else:
                buf = param_state['momentum_buffer']
                buf.mul_(momentum).add_(1 - dampening, d_p)
            if nesterov:
                d_p = d_p.add(momentum, buf)
            else:
                d_p = buf

        return -group['lr'] * d_p

```

```

[7]: from torchvision import datasets, transforms
    from torchvision.transforms import ToTensor

    image_size = 64
    transform = transforms.Compose(
        [transforms.ToTensor(),
         transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

    data_transform = transforms.Compose(
        [
            transforms.Resize(image_size),
            transforms.CenterCrop(image_size),

```

```

        transforms.ToTensor(),
        transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
    ]
)

train_data = datasets.CIFAR10(root='./data', train=True, transform =_
    ↳data_transform, download=True)

test_data = datasets.CIFAR10(root='./data', train=False, transform =_
    ↳data_transform, download=True)

print(train_data)
print(test_data)

```

Files already downloaded and verified

Files already downloaded and verified

Dataset CIFAR10

Number of datapoints: 50000

Root location: ./data

Split: Train

StandardTransform

Transform: Compose(

Resize(size=64, interpolation=bilinear, max_size=None,
antialias=None)

CenterCrop(size=(64, 64))

ToTensor()

Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))

)

Dataset CIFAR10

Number of datapoints: 10000

Root location: ./data

Split: Test

StandardTransform

Transform: Compose(

Resize(size=64, interpolation=bilinear, max_size=None,
antialias=None)

CenterCrop(size=(64, 64))

ToTensor()


Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))

)

```

[8]: import matplotlib.pyplot as plt
plt.imshow(train_data.data[0])
plt.title('%i' % train_data.targets[1])
plt.show()

```



```
[9]: batch_size = 128

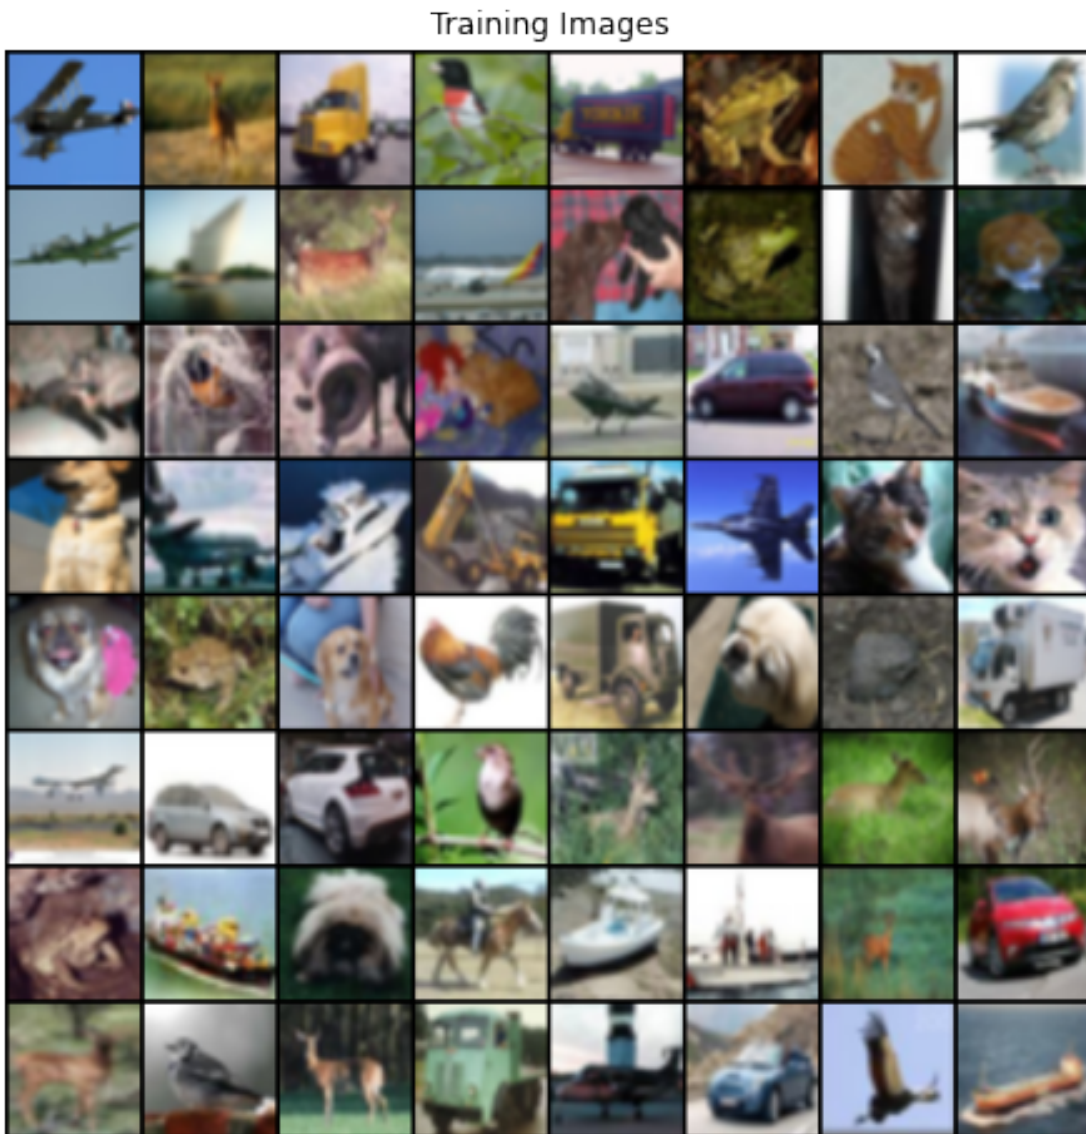
train_dataloader = idist.auto_dataloader(
    train_data,
    batch_size=batch_size,
    num_workers=2,
    shuffle=True,
    drop_last=True,
)

test_dataloader = idist.auto_dataloader(
    test_data,
    batch_size=batch_size,
    num_workers=2,
    shuffle=False,
    drop_last=True,
)
```

```
[10]: batch_size = 128
real_batch = next(iter(train_dataloader))

plt.figure(figsize=(8,8))
plt.axis("off")
plt.title("Training Images")
```

```
plt.imshow(np.transpose(vutils.make_grid(real_batch[0][:64], padding=2,
↪normalize=True).cpu(),(1,2,0)))
plt.show()
```



0.3 Models for GAN

0.3.1 Generator

The latent space dimension of input vectors for the generator is a key parameter of GAN.

```
[11]: latent_dim = 100
```

```
[12]: class Generator3x64x64(nn.Module):
    def __init__(self, latent_dim):
        super(Generator3x64x64, self).__init__()
        self.model = nn.Sequential(
            nn.ConvTranspose2d(latent_dim, 512, 4, 1, 0, bias=False),
            nn.BatchNorm2d(512),
            nn.ReLU(True),
            # state size. 512 x 4 x 4
            nn.ConvTranspose2d(512, 256, 4, 2, 1, bias=False),
            nn.BatchNorm2d(256),
            nn.ReLU(True),
            # state size. 256 x 8 x 8
            nn.ConvTranspose2d(256, 128, 4, 2, 1, bias=False),
            nn.BatchNorm2d(128),
            nn.ReLU(True),
            # state size. 128 x 16 x 16
            nn.ConvTranspose2d(128, 64, 4, 2, 1, bias=False),
            nn.BatchNorm2d(64),
            nn.ReLU(True),
            # state size. 64 x 32 x 32
            nn.ConvTranspose2d(64, 3, 4, 2, 1, bias=False),
            nn.Tanh()
            # final state size. 3 x 64 x 64
        )

    def forward(self, x):
        x = self.model(x)
        return x
```

As for dataloading, distributed models requires some specifics that `idist` addresses providing the `auto_model` helper.

```
[13]: netG = idist.auto_model(Generator3x64x64(latent_dim))
```

Note that the model is automatically moved to the best device detected by `idist`.

```
[14]: idist.device()
```

```
[14]: device(type='cuda')
```

```
[15]: summary(netG, (latent_dim, 1, 1))
```

Layer (type)	Output Shape	Param #
ConvTranspose2d-1	[-1, 512, 4, 4]	819,200
BatchNorm2d-2	[-1, 512, 4, 4]	1,024
ReLU-3	[-1, 512, 4, 4]	0
ConvTranspose2d-4	[-1, 256, 8, 8]	2,097,152

BatchNorm2d-5	[-1, 256, 8, 8]	512
ReLU-6	[-1, 256, 8, 8]	0
ConvTranspose2d-7	[-1, 128, 16, 16]	524,288
BatchNorm2d-8	[-1, 128, 16, 16]	256
ReLU-9	[-1, 128, 16, 16]	0
ConvTranspose2d-10	[-1, 64, 32, 32]	131,072
BatchNorm2d-11	[-1, 64, 32, 32]	128
ReLU-12	[-1, 64, 32, 32]	0
ConvTranspose2d-13	[-1, 3, 64, 64]	3,072
Tanh-14	[-1, 3, 64, 64]	0

Total params: 3,576,704
 Trainable params: 3,576,704
 Non-trainable params: 0

Input size (MB): 0.00
 Forward/backward pass size (MB): 3.00
 Params size (MB): 13.64
 Estimated Total Size (MB): 16.64

0.3.2 Discriminator

```
[16]: class Discriminator3x64x64(nn.Module):
    def __init__(self):
        super(Discriminator3x64x64, self).__init__()
        self.model = nn.Sequential(
            # input is 3 x 64 x 64
            nn.Conv2d(3, 64, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. 64 x 32 x 32
            nn.Conv2d(64, 128, 4, 2, 1, bias=False),
            nn.BatchNorm2d(128),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. 128 x 16 x 16
            nn.Conv2d(128, 256, 4, 2, 1, bias=False),
            nn.BatchNorm2d(256),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. 256 x 8 x 8
            nn.Conv2d(256, 512, 4, 2, 1, bias=False),
            nn.BatchNorm2d(512),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. 512 x 4 x 4
            nn.Conv2d(512, 1, 4, 1, 0, bias=False),
            nn.Sigmoid()
        )
```

```
def forward(self, x):
    x = self.model(x)
    return x
```

```
[17]: netD = idist.auto_model(Discriminator3x64x64())
summary(netD, (3, 64, 64))
```

```
-----
Layer (type)           Output Shape          Param #
=====
      Conv2d-1          [-1, 64, 32, 32]       3,072
    LeakyReLU-2          [-1, 64, 32, 32]         0
      Conv2d-3          [-1, 128, 16, 16]    131,072
    BatchNorm2d-4        [-1, 128, 16, 16]     256
    LeakyReLU-5          [-1, 128, 16, 16]         0
      Conv2d-6          [-1, 256, 8, 8]     524,288
    BatchNorm2d-7        [-1, 256, 8, 8]      512
    LeakyReLU-8          [-1, 256, 8, 8]         0
      Conv2d-9          [-1, 512, 4, 4]    2,097,152
    BatchNorm2d-10        [-1, 512, 4, 4]     1,024
    LeakyReLU-11          [-1, 512, 4, 4]         0
      Conv2d-12          [-1, 1, 1, 1]       8,192
      Sigmoid-13         [-1, 1, 1, 1]         0
=====
Total params: 2,765,568
Trainable params: 2,765,568
Non-trainable params: 0
-----

Input size (MB): 0.05
Forward/backward pass size (MB): 2.31
Params size (MB): 10.55
Estimated Total Size (MB): 12.91
-----
```

```
[18]: criterion = nn.BCELoss()
```

A batch of 64 fixed samples will be used for generating images from throughout the training. This will allow a qualitative evaluation throughout the training progress.

```
[19]: fixed_noise = torch.randn(64, latent_dim, 1, 1, device=idist.device())
```

Finally, two separate optimizers are set up, one for the generator, and one for the discriminator. Yet, another helper method `auto_optim` provided by `idist` will help to adapt optimizer for distributed configurations.

```
[20]: optimizerD = idist.auto_optim(
    ExtraSGD(netD.parameters(), lr=0.0002)
)
```

```
optimizerG = idist.auto_optim(
    ExtraSGD(netG.parameters(), lr=0.0002)
)
```

```
[21]: real_label = 1
fake_label = 0

def training_step(engine, data):
    # Set the models for training
    netG.train()
    netD.train()

    #####
    # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
    #####
    ## Train with all-real batch
    netD.zero_grad()
    # Format batch
    real = data[0].to(idist.device())
    b_size = real.size(0)
    label = torch.full((b_size,), real_label, dtype=torch.float, device=idist.
    ↪device())

    # Forward pass real batch through D
    output1 = netD(real).view(-1)
    # Calculate loss on all-real batch
    errD_real = criterion(output1, label)
    # Calculate gradients for D in backward pass
    errD_real.backward()

    ## Train with all-fake batch
    # Generate batch of latent vectors
    noise = torch.randn(b_size, latent_dim, 1, 1, device=idist.device())
    # Generate fake image batch with G
    fake = netG(noise)
    label.fill_(fake_label)
    # Classify all fake batch with D
    output2 = netD(fake.detach()).view(-1)
    # Calculate D's loss on the all-fake batch
    errD_fake = criterion(output2, label)
    # Calculate the gradients for this batch, accumulated (summed) with
    ↪previous gradients
    errD_fake.backward()
    # Compute error of D as sum over the fake and the real batches
    errD = errD_real + errD_fake
    optimizerD.extrapolation()
    # Update D
```

```

optimizerD.step()

#####
# (2) Update G network: maximize log(D(G(z)))
#####
netG.zero_grad()
label.fill_(real_label) # fake labels are real for generator cost
# Since we just updated D, perform another forward pass of all-fake batch
→ through D
output3 = netD(fake).view(-1)
# Calculate G's loss based on this output
errG = criterion(output3, label)
# Calculate gradients for G
errG.backward()
optimizerG.extrapolation()
# Update G
optimizerG.step()

return {
    "Loss_G" : errG.item(),
    "Loss_D" : errD.item(),
    "D_x": output1.mean().item(),
    "D_G_z1": output2.mean().item(),
    "D_G_z2": output3.mean().item(),
}

```

```
[22]: trainer = Engine(training_step)
```

```
[23]: def initialize_fn(m):
    classname = m.__class__.__name__
    if classname.find('Conv') != -1:
        nn.init.normal_(m.weight.data, 0.0, 0.02)
    elif classname.find('BatchNorm') != -1:
        nn.init.normal_(m.weight.data, 1.0, 0.02)
        nn.init.constant_(m.bias.data, 0)

```

```
[24]: @trainer.on(Events.STARTED)
def init_weights():
    netD.apply(initialize_fn)
    netG.apply(initialize_fn)

```

```
[25]: G_losses = []
D_losses = []

@trainer.on(Events.ITERATION_COMPLETED)
def store_losses(engine):

```

```

o = engine.state.output
G_losses.append(o["Loss_G"])
D_losses.append(o["Loss_D"])

```

```

[26]: img_list = []

@trainer.on(Events.ITERATION_COMPLETED(every=500))
def store_images(engine):
    with torch.no_grad():
        fake = netG(fixed_noise).cpu()
        img_list.append(fake)

```

```

[27]: from ignite.metrics import FID, InceptionScore

```

```

[28]: fid_metric = FID(device=idist.device())

```

```

[29]: is_metric = InceptionScore(device=idist.device(), output_transform=lambda x:
    ↪x[0])

```

```

[30]: import PIL.Image as Image

def interpolate(batch):
    arr = []
    for img in batch:
        pil_img = transforms.ToPILImage()(img)
        resized_img = pil_img.resize((299,299), Image.BILINEAR)
        arr.append(transforms.ToTensor()(resized_img))
    return torch.stack(arr)

def evaluation_step(engine, batch):
    with torch.no_grad():
        noise = torch.randn(batch_size, latent_dim, 1, 1, device=idist.device())
        netG.eval()
        fake_batch = netG(noise)
        fake = interpolate(fake_batch)
        real = interpolate(batch[0])
        return fake, real

```

```

[31]: evaluator = Engine(evaluation_step)
fid_metric.attach(evaluator, "fid")
is_metric.attach(evaluator, "is")

```

```

[32]: fid_values = []
is_values = []

```

```

@trainer.on(Events.EPOCH_COMPLETED)
def log_training_results(engine):
    evaluator.run(test_dataloader,max_epochs=1)
    metrics = evaluator.state.metrics
    fid_score = metrics['fid']
    is_score = metrics['is']
    fid_values.append(fid_score)
    is_values.append(is_score)
    print(f"Epoch [{engine.state.epoch}/5] Metric Scores")
    print(f"*   FID : {fid_score:4f}")
    print(f"*   IS : {is_score:4f}")

```

```
[33]: from ignite.metrics import RunningAverage
```

```

RunningAverage(output_transform=lambda x: x["Loss_G"]).attach(trainer, 'Loss_G')
RunningAverage(output_transform=lambda x: x["Loss_D"]).attach(trainer, 'Loss_D')

```

```
[34]: from ignite.contrib.handlers import ProgressBar
```

```

ProgressBar().attach(trainer, metric_names=['Loss_G','Loss_D'])
ProgressBar().attach(evaluator)

```

```
[35]: def training(*args):
        trainer.run(train_dataloader, max_epochs=5)
```

```
[36]: with idist.Parallel(backend='nccl') as parallel:
        parallel.run(training)
```

```
[1/390]  0%|          [00:00<?]
```

```
[1/78]   1%|1         [00:00<?]
```

```
Epoch [1/5] Metric Scores
```

```
*   FID : 0.154986
```

```
*   IS : 1.032711
```

```
[1/390]  0%|          [00:00<?]
```

```
[1/78]   1%|1         [00:00<?]
```

```
Epoch [2/5] Metric Scores
```

```
*   FID : 0.156611
```

```
*   IS : 1.037504
```

```
[1/390]  0%|          [00:00<?]
```

```
[1/78]   1%|1         [00:00<?]
```

Epoch [3/5] Metric Scores

* FID : 0.156914

* IS : 1.028521

[1/390] 0%| [00:00<?]

[1/78] 1%|1 [00:00<?]

Epoch [4/5] Metric Scores

* FID : 0.163128

* IS : 1.043960

[1/390] 0%| [00:00<?]

[1/78] 1%|1 [00:00<?]

Epoch [5/5] Metric Scores


* FID : 0.159262

* IS : 1.019424

```
[37]: %matplotlib inline
```

```
plt.figure(figsize=(10,5))
plt.title("Generator and Discriminator Loss During Training")
plt.plot(G_losses,label="G")
plt.plot(D_losses,label="D")
plt.xlabel("iterations")
plt.ylabel("Loss")
plt.legend()
```

```
[37]: <matplotlib.legend.Legend at 0x7fcaa37d6ca0>
```



```
[38]: fig, ax1 = plt.subplots()


plt.title("Evaluation Metric During Training")

color = 'tab:red'
ax1.set_xlabel('epochs')
ax1.set_ylabel('IS', color=color)
ax1.plot(is_values, color=color)

ax2 = ax1.twinx()

color = 'tab:blue'
ax2.set_ylabel('FID', color=color)
ax2.plot(fid_values, color=color)

fig.tight_layout()
```



```
[39]: %matplotlib inline

# Grab a batch of real images from the dataloader
real_batch = next(iter(train_dataloader))
```



```

# Plot the real images
plt.figure(figsize=(15,15))
plt.subplot(1,2,1)
plt.axis("off")
plt.title("Real Images")
plt.imshow(np.transpose(vutils.make_grid(real_batch[0][:64], padding=5,
↳normalize=True).cpu(), (1,2,0)))

# Plot the fake images from the last epoch
plt.subplot(1,2,2)
plt.axis("off")
plt.title("Fake Images")
plt.imshow(np.transpose(vutils.make_grid(img_list[-1], padding=2,
↳normalize=True).cpu(), (1,2,0)))

```

[39]: <matplotlib.image.AxesImage at 0x7fcaccd56460>

